

Query Optimization Techniques



@data insights

What is Query Optimization?

- **Process of enhancing database queries for faster execution.**
- **Reduces resource usage (CPU, memory, I/O).**
- **Improves user experience and system scalability.**

Query




Optimization



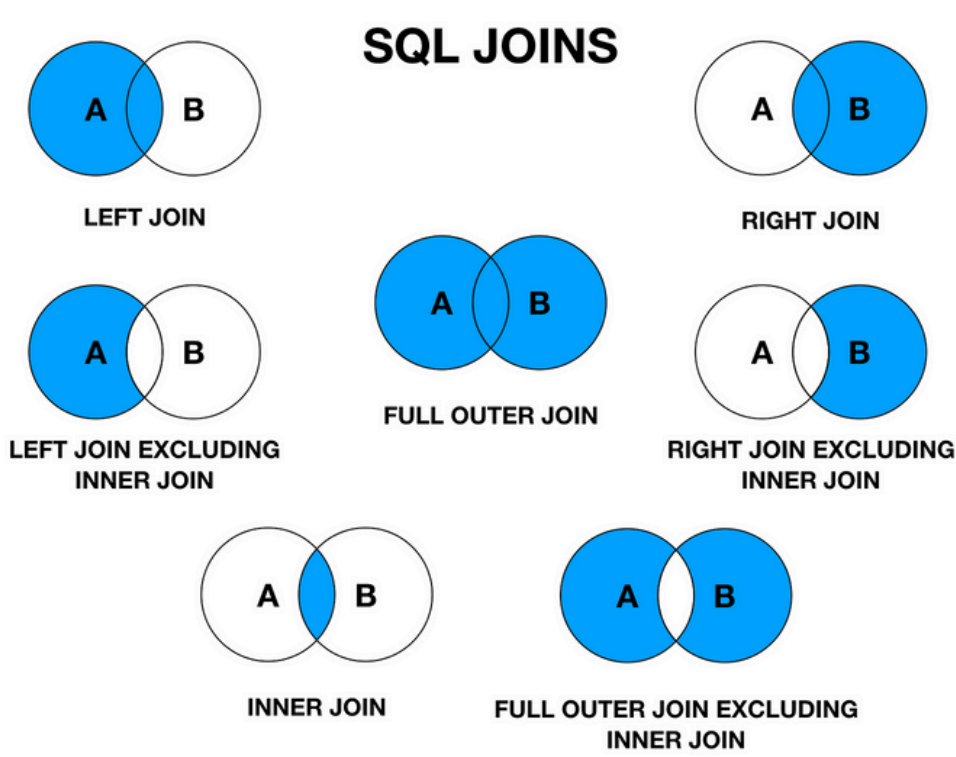
Execution

Common Query Optimization Techniques

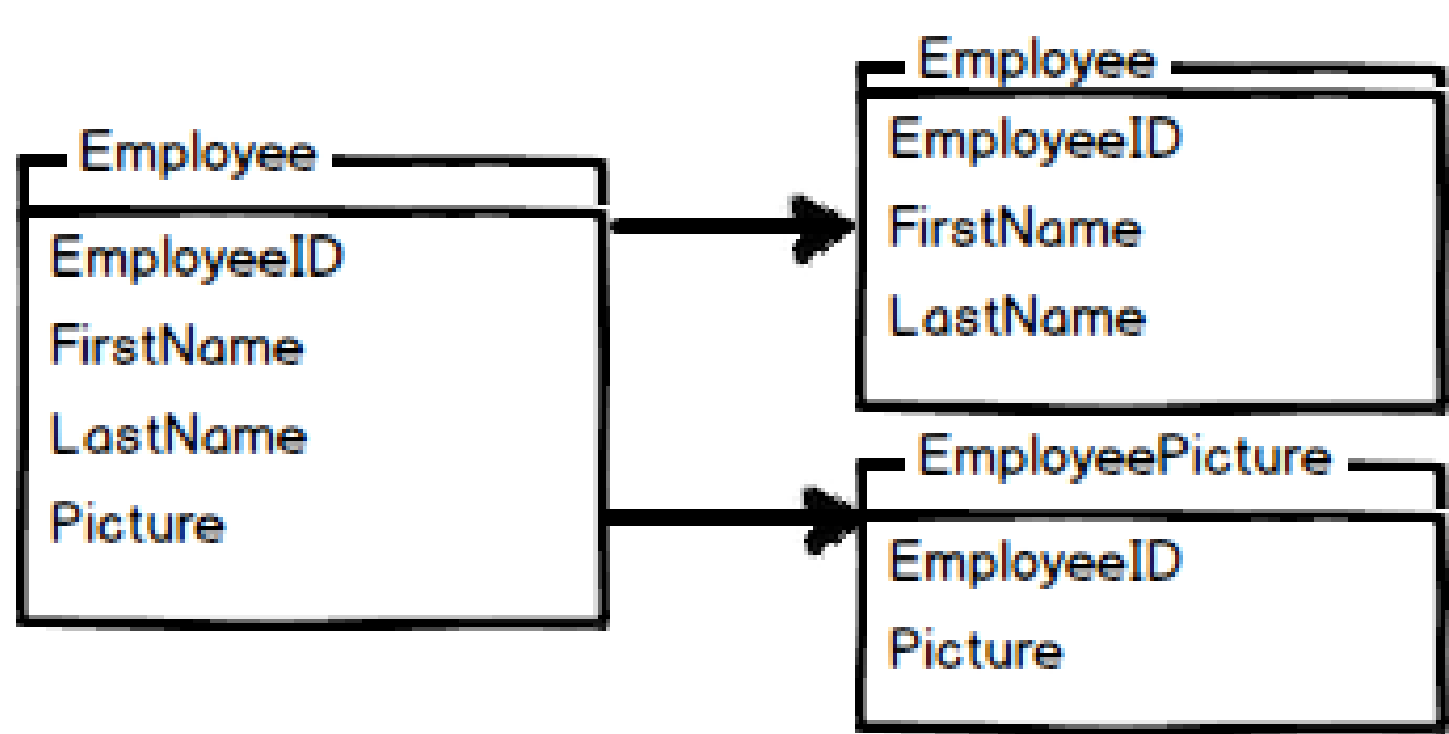
- Indexing



- Query Refactoring
- Avoiding SELECT *
- Using Joins Instead of Subqueries



- Partitioning Large Tables



Before/After Example 1 - Indexing

Before:

- Query: `SELECT * FROM users WHERE last_name = 'Smith';`
- Execution Plan: Full Table Scan.
- Performance: Slow (e.g., 500ms).

id	first_name	last_name	email	age
1	John	Smith	john@example.com	30
2	Alice	Johnson	alice@example.com	25
3	Bob	Smith	bob@example.com	40
4	Emily	Davis	emily@example.com	35
5	Mark	Brown	mark@example.com	28

Since there is no index on `last_name`, the database performs a Full Table Scan—checking each row one by one.

After:

- Query: `CREATE INDEX idx_last_name ON users(last_name);`
- Execution Plan: Index Seek.
- Performance: Fast (e.g., 50ms).

id	first_name	last_name	email	age
1	John	Smith	john@example.com	30
3	Bob	Smith	bob@example.com	40
2	Alice	Johnson	alice@example.com	25
4	Emily	Davis	emily@example.com	35
5	Mark	Brown	mark@example.com	28

Since an index exists on `last_name`, the database can directly find matching rows.

Before/After Example 2 - Avoiding SELECT

Before:

- Query: `SELECT * FROM orders;`
- Execution Plan: Scans all columns.
- Performance: Slow (e.g., 300ms).

After:


- Query: `SELECT "Order ID", "Customer Name" FROM orders;`
- Execution Plan: Scans only required columns.
- Performance: Fast (e.g., 100ms).

Order ID	Customer Name	Product	Quantity	Price	Order Date	Status
1001	John Doe	Laptop	1	\$999	2025-02-07	Shipped
1002	Jane Smith	Phone	2	\$799	2025-02-06	Pending
1003	Bob Brown	Tablet	1	\$499	2025-02-05	Delivered
1004	Alice Green	Monitor	3	\$299	2025-02-04	Shipped
1005	Michael White	Keyboard	5	\$49	2025-02-03	Delivered
1006	Sarah Black	Headphones	2	\$199	2025-02-02	Pending
1007	Tom Wilson	Mouse	4	\$29	2025-02-01	Shipped
1008	Emma Johnson	Smartwatch	1	\$249	2025-01-31	Delivered
1009	David Clark	Printer	1	\$199	2025-01-30	Canceled
1010	Olivia Adams	Speakers	2	\$149	2025-01-29	Shipped
1011	Liam Martin	Laptop	1	\$1099	2025-01-28	Pending
1012	Noah Carter	Tablet	2	\$599	2025-01-27	Delivered
1013	Sophia Evans	Phone	1	\$899	2025-01-26	Canceled
1014	James Scott	Gaming PC	1	\$1499	2025-01-25	Shipped
1015	Ava Wright	Monitor	↓	\$399	2025-01-24	Pending

Before/After Example 3 - Using Joins Instead of Subqueries

Before:

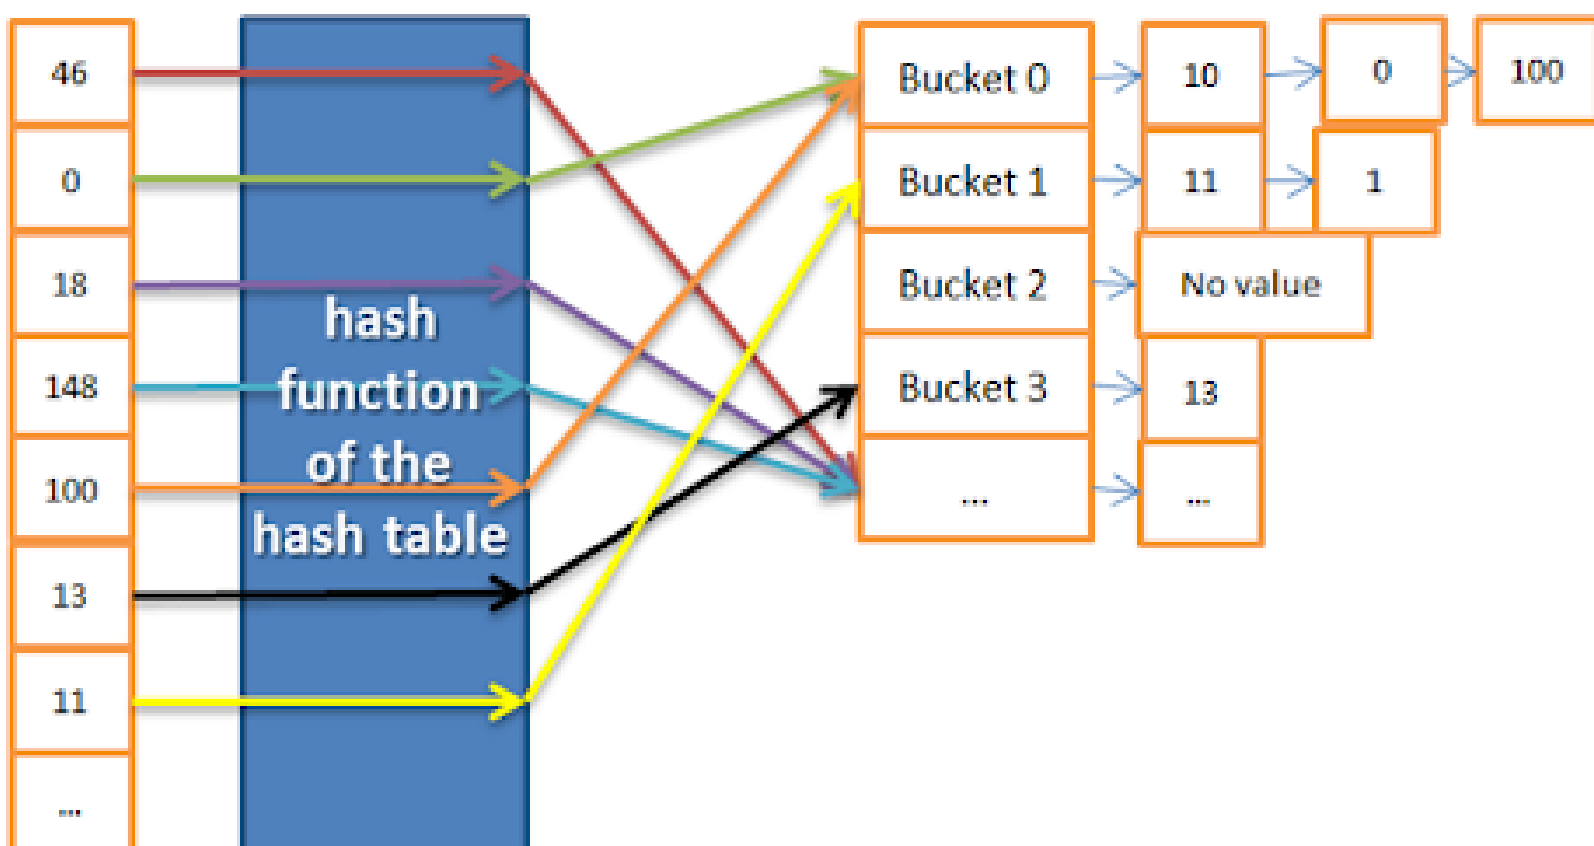
- Query: `SELECT * FROM users WHERE user_id IN (SELECT user_id FROM orders);`
- Execution Plan: Nested Loop.
- Performance: Slow (e.g., 400ms).



After:

- Query: `SELECT u.* FROM users u JOIN orders o ON u.user_id = o.user_id;`
- Execution Plan: Hash Join.
- Performance: Fast (e.g., 150ms).

Hash Join



Execution Plans Explained

An execution plan is a roadmap of how a database query is executed, showing the steps the optimizer chooses to retrieve data efficiently.

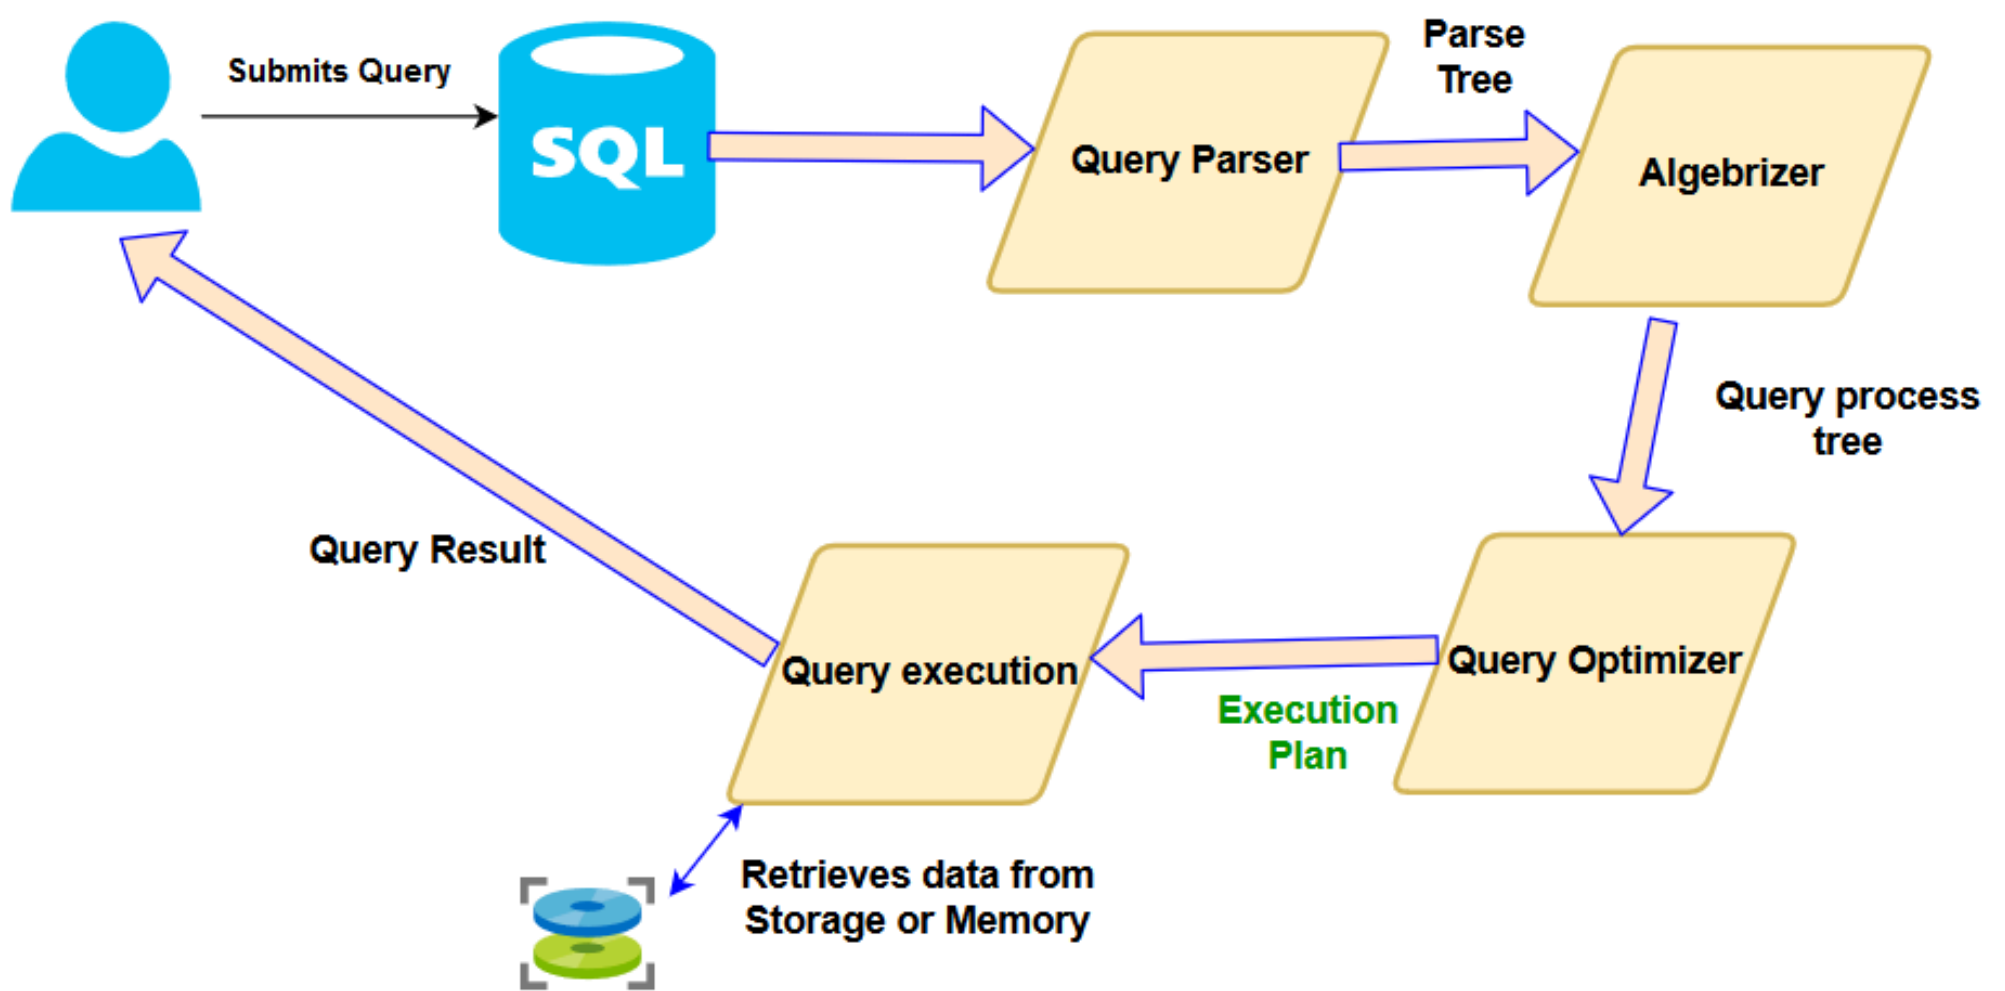
How to Read It:

Operations:

- **Scan (Full table/index scan)** → Reads all rows, less efficient.
- **Seek (Index seek)** → Quickly finds specific rows using an index.
- **Join (Nested Loop, Hash, Merge)** → Combines data from multiple tables.

Cost: Represents the estimated resource usage; lower is better.

Rows: Estimated number of rows processed at each step.



Conclusion

Key Points:

- **Query Optimization is Critical:** Optimized queries reduce execution time and resource consumption.
- **Small Changes, Big Impact:** Indexing, rewriting queries, or restructuring joins can dramatically improve performance.
- **Measure & Test:** Always analyze execution plans and benchmark performance before and after changes.

Like
Comment
Share



Data **INSIGHTS**